

OMISSION

THE OMISSION STANDARD

A Triangulated Epistemological Framework for Child Welfare Accountability

Version 2.2

DOCUMENT STATUS: Active Standard
PRIMARY AUDIENCES: Regulators · Legal Practitioners · Independent Auditors · Academic Researchers
PLATFORM ALIGNMENT: Fully aligned with OMISSION Platform (sbouel72.github.io/omission-platform)

PURPOSE	This Standard provides: (I) the theoretical justification for re-weighting evidentiary sources in institutional child welfare assessments; (II) the eight survivor-centred domain definitions and their operational proxy indicators; (III) the scoring model aligned with the OMISSION Platform H×W architecture; (IV) the cross-cutting audit infrastructure (Process, Compliance, Oversight) that applies across all domains; (V) the evidentiary weighting rules for adjudicating conflicting claims; and (VI) the versioning and review protocol.
v2.1 → v2.2	Domain taxonomy replaced: the eight institutional accountability domains are superseded by the eight survivor-centred rights domains from the OMISSION Platform. Process, Compliance, and Oversight are reclassified as cross-cutting audit infrastructure. Scoring outputs renamed and aligned to the Platform H×W model.

Prepared by the OMISSION Platform
sbouel72.github.io/omission-platform

PART I: THEORETICAL FOUNDATION

1.1 The Asymmetry Problem

In conventional child welfare assessments, administrative archives and clinical case files are accorded presumptive evidentiary weight. This default weighting is structurally inherited from the institutional systems that produce them — creating a closed loop in which the system assesses itself and treats its self-assessment as objective truth.

The OMISSION Standard does not propose the wholesale rejection of clinical or administrative data. It proposes a triangulated weighting heuristic that distinguishes between two qualitatively different categories of secondary evidence, and positions first-person survivor testimony as the epistemological anchor against which all other evidence must be triangulated.

1.2 Evidence Categories

Category A: Corroborative Clinical Literature

Studies, reports, or assessments satisfying all of the following criteria carry contextual corroborative weight:

- Methodological transparency: data collection methods, sample bias, funding sources, and conflicts of interest are fully and proactively disclosed.
- Genuine informed consent: subjects were fully informed, voluntarily participated, and their right to withdraw without penalty was explicitly documented.
- Replicability and independent verification: the methodology and underlying data are accessible for independent audit and re-analysis.
- Institutional independence: the authors or commissioning body do not carry a direct institutional interest in the outcome of the specific case being assessed.

Category B: Institutional Self-Reports

Assessments, agency publications, clinical notes, or internal investigations exhibiting one or more of the following characteristics are treated as data about the institution's self-presentation — not as neutral accounts of what occurred:

- Originate from the same institution or a body with direct institutional interest in the case being assessed.
- Lack transparency around methodology, funding, or data provenance.
- Are not independently verifiable due to proprietary claims or confidentiality assertions.
- Function as exculpatory documentation — produced in response to an allegation or complaint and designed to defend the institution's position.

LEGAL PARALLEL

This distinction is structurally identical to how legal systems treat self-serving hearsay versus direct testimony: the probative value of a statement is diminished when its

author has a direct interest in the conclusion it supports. Category B evidence is not disqualified — it is contextualised.

1.3 The Primacy of Lived Archive

First-person survivor testimony occupies a distinct epistemic category. It is not 'more objective' than clinical or institutional data; it is ontologically different. It is the only evidentiary category that emerges directly from the experience of the system's effects. As such, it functions as the anchor point against which all other data must be triangulated.

When Category B evidence contradicts survivor testimony, the burden of explanation falls on the institution to demonstrate — through Category A criteria and independently verifiable documentary evidence — why its account should be preferred. Where no such evidence exists and the institution's Visibility Deficit is elevated, the default epistemological position favours the survivor's account.

1.4 The Eight Survivor-Centred Domains

Version 2.2 adopts the domain taxonomy of the OMISSION Platform. The eight domains are survivor-centred rights domains — they name what is actually at stake for the people the system acts upon. They replace the institutional accountability domains used in v2.0 and v2.1, which are reclassified as cross-cutting audit infrastructure (see Section 1.5).

Domain	Operational Definition	What Institutional Failure Looks Like
1. Identity Continuity	Whether people can know and maintain a consistent record of who they are: name, origin, genealogy, and legal identity across the full arc of their involvement with the system.	Identity records altered, sealed, or denied without consent; birth names replaced without documentation; access to original birth certificates blocked; identity treated as institutional property rather than personal right.
2. Family Integrity	Whether family relationships and connections are preserved and respected: contact, knowledge of family members, and the right to maintain bonds across separation.	Family contact terminated without independent review; siblings separated without documented justification; family members denied information about a child's whereabouts or wellbeing; reunification systematically deprioritised.
3. Cultural Continuity	Whether people retain access to their cultural, linguistic, and community heritage: language, customs, community connections, and cultural identity.	Children placed in culturally incompatible environments without documented consideration; language of origin lost without record; cultural heritage treated as secondary to administrative convenience; no cultural continuity planning in placement decisions.
4. Information Rights	Whether people can access records and information that directly concern them: case files, decisions made about them, and	Records sealed beyond statutory requirements; FOI requests delayed or refused without specific documented exemption; case files contain third-party paraphrases of the subject's own words

Domain	Operational Definition	What Institutional Failure Looks Like
	the reasoning behind those decisions.	without primary source; destruction of records without survivor consent.
5. Agency	Whether affected people can act meaningfully on their own behalf within the system: participate in decisions, have their preferences documented, and receive responses to their representations.	Survivor representations unacknowledged or undocumented; decisions made without survivor participation at any stage; institutional 'best interest' determinations made without survivor input; complaint mechanisms inaccessible or effectively non-functional.
6. Consent	Whether decisions affecting people are made with their informed, voluntary, and ongoing agreement: documented consent processes, right to withdraw, and freedom from coercion.	Consent obtained under duress or without full information; consent to one thing used to authorise another; no documented consent process; withdrawal of consent not respected; children and young people excluded from consent processes affecting them.
7. Self-Determination	Whether affected groups can meaningfully shape the rules, policies, and governance structures that govern them: representation, consultation, and co-design.	Policy developed without survivor representation; consultation processes tokenistic or post-decisional; survivor advisory bodies lack decision-making power; systemic feedback loops absent; affected communities excluded from framework design.
8. Memory Preservation	Whether the historical record of what happened to people is kept accurately, completely, and accessibly over time: retention, correction, and survivor access to their own history.	Records destroyed, lost, or degraded without survivor consent; corrections to inaccurate records refused or undocumented; transition between agencies results in record loss; survivors denied copies of their own files; historical records rewritten to serve institutional interests.

1.5 Cross-Cutting Audit Infrastructure

Process, Compliance, and Oversight are not domains in v2.2. They are audit infrastructure that applies across all eight domains. For every domain assessment, the auditor must apply all three infrastructure layers. This structure ensures that the procedural accountability questions are asked of every substantive rights domain, rather than occupying a parallel track that can be scored independently of survivor-centred outcomes.

Infrastructure Layer	Function	How It Applies Across Domains
Process	Documents the sequence, rationale, and decision-maker accountability for every institutional action across all eight domains.	For each domain: is there a documented decision trail? Are decision-makers identified and their credentials recorded? Is the rationale for each key decision explicit and independently reviewable?
Compliance	Measures adherence to statutory obligations, regulatory requirements,	For each domain: what are the applicable statutory and regulatory obligations? Are they

Infrastructure Layer	Function	How It Applies Across Domains
	and internal policy — in letter and in spirit — across all eight domains.	met? Where exceptions occur, are they documented and justified? What do external audits reveal vs. internal self-assessment?
Oversight	Assesses whether independent accountability structures exist and function across all eight domains: who watches the watchers, and whether their findings are accessible.	For each domain: is there an independent oversight body? How frequently does it audit? Are its findings accessible to survivors and the public? Does it have genuine remedial power?

PART II: SCORING MODEL

2.1 Alignment with the OMISSION Platform

Version 2.2 aligns the Standard's scoring outputs directly with the OMISSION Platform's $H \times W$ model. A complete Standard audit now produces scores that can be imported directly into the Platform's domain grid, enabling the Platform's radar visualisation, gap analysis, and reference signature comparison to be populated from auditor-verified evidence rather than self-assessment alone.

The Platform's two assessment modes — Quick Assessment (self-reported H and W per domain) and Deep Scan (AI-analysed document scoring) — are both enhanced by a Standard audit: auditor-derived H and W scores per domain replace or validate self-reported or AI-generated scores, producing a fully evidenced Platform output.

Output	Definition	Platform Alignment
Visibility Deficit (H)	Per-domain score (0–1) measuring the harm potential of institutional opacity in that domain. Derived from the proxy indicator aggregate. High score = high harm potential from what the institution is not disclosing or not seeing.	Maps directly to H (Harm potential) in the Platform's Quick Assessment and Deep Scan modes. A Standard audit produces an H score per domain that can be imported into the Platform.
Disclosure Score (W)	Per-domain score (0.01–1) measuring how well the institution actually discloses and makes visible the risks and information in that domain. Derived from proxy indicator thresholds.	Maps directly to W (Visibility weight) in the Platform. A Standard audit produces a W score per domain that can be imported into the Platform.
Omission Risk Index	Per-domain score: $H \times (1-W)$. The gap between harm potential and actual visibility. High score = high risk hiding in institutional blind spots.	Identical to the Platform's Omission Risk Index formula. A fully audited institution produces eight Omission Risk Index scores — one per domain — that populate the Platform's radar directly.
Domain Score	Overall domain health score (0–100): $100 - (\text{Omission Risk Index} \times 100)$. Higher = better governance. Used in the Platform's Current Score and True Score outputs.	Feeds the Platform's Current Score (perceived governance health) vs True Score (actual harm potential) gap analysis.

2.2 Per-Domain Scoring Methodology

Each of the eight domains is scored independently. The domain score is derived from the five proxy indicators (Section 2.3) applied to the documentation, records, and testimony relevant to that domain. Cross-cutting audit infrastructure (Process, Compliance, Oversight) is applied within each domain assessment rather than scored separately.

SCORING PRINCIPLE

Scores reflect documented evidence only. An institution's failure to produce documentation is itself scored — as an indicator of opacity, not as a neutral absence. The absence of evidence is not the evidence of absence; in institutional accountability contexts, it is frequently the evidence of concealment. A domain with no documentation is scored at maximum Visibility Deficit, not excluded from the assessment.

2.3 Proxy Indicators and Deficit Thresholds

The five proxy indicators are applied per domain. Each indicator is assessed against its threshold; triggered thresholds contribute to the domain's Visibility Deficit (H) score and reduce the Disclosure Score (W). All five proxy indicators apply to all eight domains — there are no domain-specific exemptions from the indicator set, though the specific documentary evidence examined will vary by domain.

Indicator	Definition	Measurement	Deficit Threshold & Derivation
MMDR Minimum Mandatory Disclosure Ratio	Ratio of legally required disclosure items actually disclosed, voluntarily or on request, against the full statutory and regulatory disclosure schedule for the relevant domain.	Required disclosure items (from applicable statute/regulation per domain) vs. actual disclosure items in the case file. Assessed per domain.	Ratio < 0.80 flags a deficit. Derivation: benchmarked against public sector disclosure compliance literature establishing 80% as the minimum credible voluntary compliance threshold. Per-domain statutory schedules are documented in the OMISSION Platform domain reference files.
FOIRL FOI / Records Response Latency	Time taken to respond to formal information or records access requests relating to any of the eight domains.	Average calendar days to response vs. statutory maximum for the relevant jurisdiction. Assessed per domain where distinct request categories apply.	Latency > 1.5× the statutory maximum flags a deficit. Multiple consecutive breaches within 12 months elevate the severity classification.
RD Redaction Density	Proportion of text redacted in released documents, assessed per domain across a minimum sample of 10 documents (or the full file where fewer exist).	Percentage of redacted words or lines per released document, per domain. Specific attention to redaction of survivor-identifying information vs. third-party information.	RD > 30% (absent a specific, documented, and independently reviewable statutory exemption) flags a deficit. Where redaction disproportionately affects survivor-generated content vs. institutional content, severity is elevated.
NIS	Frequency and severity of documented contradictions between internal documents	Independent comparative audit of internal and external	NIS > 3 Significant or Critical contradictions per 100 pages flags a deficit.

Indicator	Definition	Measurement	Deficit Threshold & Derivation
Narrative Inconsistency Score	and public-facing statements, classified as Minor, Significant, or Critical per domain.	document sets per domain. Minor: factual discrepancy without material consequence. Significant: discrepancy affecting case outcome. Critical: discrepancy suggesting deliberate misrepresentation.	Derivation: adapted from PCAOB AS 2201 document consistency standards for welfare case file review. Per-domain NIS scores are weighted by domain harm potential (H).
RISA Risk Indicator Symmetry Assessment	Whether the institution's risk assessments across all eight domains include survivor-defined risk indicators, or whether risk is assessed asymmetrically in the institution's favour.	Review of all risk assessment documentation per domain. Ratio of survivor-identified risk factors included vs. excluded. Presence or absence of institutional liability risk in the same assessment.	Asymmetry ratio > 2:1 (institutional : survivor-identified) flags a deficit. Complete absence of survivor-defined risk indicators in any formal domain assessment is an automatic Critical flag. Applied across all eight domains, not only Domain 4.

2.4 The Omission Risk Index and Inversion Curve

The Omission Risk Index is calculated per domain as $H \times (1-W)$. This is identical to the Platform formula. The Inversion Curve maps each domain's Omission Risk Index to a four-tier rating, determining how institutional confidence scores are treated in that domain:

H Score	W Score	Omission Risk Index Interpretation	Auditor Action
0.00–0.25	0.75–1.00	Low risk. Institution is transparent in this domain. Institutional confidence provisionally credible subject to Category A corroboration.	Standard review. Document corroboration sources. Note any proxy indicator exceptions.
0.26–0.50	0.50–0.74	Moderate risk. Partial opacity. Institutional confidence requires Category A corroboration to be treated as credible.	Elevated scrutiny. Identify and assess all Category A corroboration. Document gaps. Flag domain in executive summary.
0.51–0.75	0.25–0.49	High risk. Active opacity. Institutional confidence treated as data about self-presentation, not objective evidence.	High scrutiny. Institutional claims not accepted without independent corroboration. Notify oversight body. Survivor testimony anchors the domain assessment.
0.76–1.00	0.01–0.24	Critical risk. Structural opacity. INVERSION APPLIED: high institutional confidence is flagged as	Critical alert. Refer to independent oversight. Institutional confidence scores inverted in the final report. Extraordinary Category A

H Score	W Score	Omission Risk Index Interpretation	Auditor Action
		a governance failure and a marker of successful obfuscation.	corroboration required before any institutional claim is credited.

2.5 Aggregate Scoring

The aggregate governance score across all eight domains is the mean of the eight Domain Scores, weighted by domain H score (harm potential). Domains with higher harm potential carry proportionally greater weight in the aggregate. This weighting ensures that opacity in high-stakes domains (Identity Continuity, Family Integrity, Memory Preservation) is not averaged away by transparency in lower-stakes domains.

The aggregate score feeds the Platform's Current Score vs. True Score gap analysis: Current Score reflects what the institution perceives itself to be doing; True Score reflects actual harm potential across domains.

PART III: EVIDENTIARY WEIGHTING RULES

3.1 Weighting Heuristic

Evidentiary weights are applied per domain. A source category's weight in any given domain is adjusted by that domain's Visibility Deficit (H) score. The same document may carry different weight in different domains depending on each domain's transparency profile.

Source Category	Baseline Weight	Adjustment Rules	Visibility Deficit Interaction
Survivor Testimony	Anchor Point (Burden-Shifting Trigger)	Not numerically weighted. Functions as the reference frame against which all other evidence is assessed. Cannot be displaced by Category B evidence alone.	A higher domain Visibility Deficit (H) score strengthens the epistemological case for the survivor's account where institutional evidence conflicts. At $H > 0.75$, survivor testimony is the primary evidentiary basis for the domain assessment.
Category Evidence A	High Corroborative Weight	Weight proportional to methodological rigour, independence, and transparency. May contextualise but not override survivor testimony.	Retains full weight regardless of the institution's Visibility Deficit score, provided it meets all four Category A criteria independently of the institution being assessed.
Category Evidence B	Low Weight (Self-Presentation Data)	Treated as data about the institution's narrative. Not objective fact.	Weight reduced by 50% where domain H is 0.26–0.50; by 75% where H is 0.51–0.75; treated as narrative data only (no evidential weight) where H is 0.76–1.00.
Non-Survivor Third-Party Testimony	Moderate Weight	Evaluated on credibility, independence, and opportunity to observe. Third parties with employment or contractual ties to the institution are reclassified as Category B.	Weight reduced in proportion to any institutional relationship the third party holds.

3.2 The Adjudication Protocol

Where evidence sources conflict within a domain, the following protocol applies in sequence. Auditors must document which rule was applied and provide a written justification for any departure from the default position.

Rule 1: Survivor Testimony vs. Category B Evidence

Survivor testimony is preferred unless the institution provides Category A evidence and independently verifiable documentary evidence that directly contradict the survivor's account on a specific, material factual claim within the domain being assessed. The institution bears the burden of producing such evidence. The auditor is not required to seek it on the institution's behalf.

Rule 2: Survivor Testimony vs. Category A Evidence

Category A evidence may contextualise or provide independent validation of specific factual claims in survivor testimony. Where it genuinely contradicts survivor testimony on a material fact, the auditor must:

1. Assess the relative methodological strengths of the Category A evidence and the conditions under which survivor testimony was recorded.
2. Audit the Category A evidence's underlying assumptions for structural biases that may disadvantage survivor accounts — including sampling populations, funding sources, and institutional affiliations.
3. Document the contradiction explicitly. Contradiction does not automatically invalidate survivor testimony; it triggers a deeper audit.
4. Where the contradiction cannot be resolved, present both accounts with their respective evidentiary strengths and limitations documented. An unresolved contradiction is a finding, not a silence.

Rule 3: Category A vs. Category B Evidence

Category A evidence supersedes Category B evidence where the two conflict, regardless of the institutional seniority or apparent authority of the Category B source.

Rule 4: High Institutional Confidence + High Visibility Deficit

Where an institution scores highly on its own confidence metrics in a domain but that domain exhibits $H > 0.50$, the institutional confidence score is inverted in the final report. High confidence is flagged as a governance risk, not an indicator of safety.

Rule 5: Conflicting Survivor Accounts

Where multiple survivor accounts conflict, each is assessed individually. Conflict between survivor accounts does not reduce the epistemological weight of survivor testimony as a category. The auditor must assess differential access to information, differential power within the institutional relationship, and the conditions under which each account was recorded.

Rule 6: Domain-Specific Burden Shift

In domains where the institution's harm potential (H) exceeds 0.75, the standard burden of proof for institutional claims is reversed: the institution must affirmatively demonstrate — through Category A evidence — that its account of events in that domain is accurate. Absence of such evidence is scored as a Critical deficit in the domain's NIS indicator.

**AUDITOR
OBLIGATION**

The adjudication protocol does not permit silence. Where evidence conflicts and a resolution cannot be reached, the auditor must document the conflict, specify the rules applied, and present the unresolved tension as a finding in the final report. An unresolved evidential conflict in a child welfare context is a governance failure — not a methodological limitation to be noted and set aside.

PART IV: VERSIONING AND REVIEW PROTOCOL

4.1 Dynamic Instrument Principle

This Standard is a dynamic instrument maintained in alignment with the OMISSION Platform. Its domain taxonomy, scoring outputs, and operational thresholds are subject to revision in response to platform updates, jurisprudential developments, empirical challenges, and operational findings. Version alignment between the Standard and the Platform is a primary maintenance obligation.

PLATFORM SYNC

Whenever the OMISSION Platform updates its domain taxonomy or H×W formula, a corresponding Standard review is triggered within 30 days. The Standard and Platform must not diverge on domain definitions, scoring outputs, or formula architecture. Any divergence is a governance failure of the Standard itself.

4.2 Review Triggers

Trigger	Specific Condition	Review Type
Scheduled	24 months from v2.2 ratification date.	Full review. All sections assessed for currency and operational adequacy.
Platform Update	A material change to the OMISSION Platform's domain taxonomy, scoring formula, or infrastructure that affects alignment with this Standard.	Alignment review. Affected sections re-calibrated and version updated within 30 days of platform release.
Jurisprudential Change	A binding legal decision materially affecting evidentiary treatment of survivor testimony or institutional self-reports in any covered jurisdiction.	Targeted amendment. Affected sections revised and reissued as a point release.
Threshold Challenge	A formally documented, evidence-based challenge demonstrating a proxy indicator threshold produces systematically erroneous results.	Threshold review. Derivation re-examined; if challenge sustained, revised with full rationale documentation.
Operational Finding	An auditor documents a case where the framework produced a finding that demonstrably failed to protect a survivor or misclassified an institutional failure.	Emergency review. Affected section suspended pending review and reissued within 90 days.

4.3 Version Numbering

- Major increment (e.g. 2.0 → 3.0): fundamental revision to the theoretical framework or evidentiary weighting architecture.
- Minor increment (e.g. 2.1 → 2.2): domain taxonomy changes, new proxy indicators, scoring formula changes, or new adjudication rules.
- Patch increment (e.g. 2.2 → 2.2.1): targeted amendment in response to a specific jurisprudential trigger or documented operational error.

4.4 Change Log

Version	Date	Summary
1.0	—	Initial release. Theoretical framework, eight institutional accountability domains, four-tier Visibility Deficit scale, basic evidentiary weighting rules.
2.0	—	Full operationalisation: proxy indicators added; domain operational definitions formalised; Adjudication Protocol introduced; MMDR, FOIRL, RD, and NIS thresholds specified.
2.1	—	NIS and MMDR threshold derivations documented. RISA added as fifth proxy indicator. Domain 4 (Risk) supplementary guidance added. Category B weight reduction schedule linked to Visibility Deficit tiers. Rule 5 (Conflicting Survivor Accounts) added. Part IV (Versioning) introduced. Third-party testimony reclassification rule added.
2.2	Current	Platform alignment release. Eight survivor-centred domains replace eight institutional accountability domains. Process, Compliance, and Oversight reclassified as cross-cutting audit infrastructure applied per domain. Scoring outputs renamed: Visibility Deficit (H), Disclosure Score (W), Omission Risk Index ($H \times (1 - W)$) — directly aligned with OMISSION Platform formula. RISA extended to all eight domains. Rule 6 (Domain-Specific Burden Shift) added. Aggregate scoring weighted by domain H score.

CLOSING PROVISION

This Standard is a structured, transparent, and auditable analytical grid through which any claim — institutional, clinical, or survivor-led — can be assessed for its epistemic provenance, its transparency, and its alignment with verifiable evidence.

The goal is not to replace clinical expertise with advocacy. The goal is to ensure that clinical expertise and institutional authority are held to the same standard of transparency and accountability as every other category of evidence. Institutions that operate with genuine integrity have nothing to fear from this Standard. Institutions that resist it, on grounds of professional authority or procedural precedent, are providing data.

Version 2.2 completes the alignment between this Standard and the OMISSION Platform. A Standard audit now produces scores that populate the Platform directly. The Platform's radar, gap analysis, and reference signature comparison are now grounded in auditor-verified evidence rather than self-assessment alone. The two instruments are a single system.

END OF THE OMISSION STANDARD, VERSION 2.2

sbouel72.github.io/omission-platform